

Heterogeneous CPUs+GPUs Computing: Models, Tools, and Applications

References

- [1] J. Fang, A.L. Varbanescu and H.J. Sips (2011). A Comprehensive Performance Comparison of CUDA and OpenCL. In The 40-th International Conference on Parallel Processing (ICPP'11), Taipei, Taiwan.
- [2] J. Shen, J. Fang, H. J. Sips, and A. L. Varbanescu. An application-centric evaluation of OpenCL on multi-core CPUs. *Parallel Computing*, vol. 39, no. 12, pp. 834 – 850, 2013.
- [3] J. Shen, A. L. Varbanescu, H. J. Sips, M. Arntzen, D. G. Simons: Glinda: A Framework for Accelerating Imbalanced Applications on Heterogeneous Platforms. In *Conf. Computing Frontiers* 2013.
- [4] J. Shen, H. J. Sips, P. Zou, Y. Lu, A. L. Varbanescu. Matching Imbalanced Workloads with Heterogeneous Platforms. *ICS* 2014.
- [5] Jie Shen and Ana Lucia Varbanescu and Henk Sips. Look Before You Leap: Using the Right Hardware Resources to Accelerate Applications. *IEEE HPCC* 2014.
- [6] Jie Shen and Ana Lucia Varbanescu and Xavier Martorell and Henk Sips (2015). A Study of Application Kernel Structure for Data Parallel Applications. Technical Report PDS-2015-001, Delft University of Technology.
- [7] Jie Shen and Ana Lucia Varbanescu and Xavier Martorell and Henk Sips (2015). Matchmaking Applications and Partitioning Strategies for Efficient Execution on Heterogeneous Platforms. *ICPP* 2015.
- [8] Jie Shen and Ana Lucia Varbanescu and Henk Sips (2015). Improving Application Performance by Efficiently Utilizing Heterogeneous Many-core Platforms. *CCGrid* 2015.
- [9] Jie Shen and Ana Lucia Varbanescu and Yutong Lu and Peng Zou and Henk Sips (2016). Workload Partitioning for Accelerating Applications on Heterogeneous Platforms. *IEEE Transactions on Parallel and Distributed Systems*. [Bibtex]
- [10] Yong Guo, Ana Lucia Varbanescu, Dick Epema and Alexandru Iosup (2016). Design and Experimental Evaluation of Distributed Heterogeneous Graph-Processing Systems. *CCGrid'16*

Other relevant references

- [11] Khronos Group., "The OpenCL Specification Version: 1.2 Document Revision 15." <http://www.khronos.org/registry/cl/specs/opencl-1.2.pdf>, November 2011.
- [12] Intel Inc., Writing Optimal OpenCL Code with Intel OpenCL SDK.
- [13] J. van der Sande, "Evaluating the Performance and Portability of OpenCL", MSc Thesis, TU Eindhoven, August 2011.
- [14] K. Karimi, N. Dickson, and F. Hamze, "A performance comparison of CUDA and OpenCL," Arxiv preprint arXiv:1005.2581, 2010.
- [15] R. Membarth, F. Hannig, J. Teich, M. K"orner, and W. Eckert, "Framework for multi-core architectures: a comprehensive evaluation using 2D/3D image registration," Architecture of Computing Systems-ARCS 2011, pp. 62–73, 2011.
- [16] C. K. Akhtar Ali, Usman Dastgeer, "OpenCL for programming shared memory multicore CPUs," in Proceedings of the 5th Workshop on MULTIPROG, in conjunction with HiPEAC 2012, January 2012.
- [17] R. Ferrer, J. Planas, P. Bellens, A. Duran, M. Gonzalez, X. Martorell, R. Badia, E. Ayguade, and J. Labarta, "Optimizing the exploitation of multicore processors and GPUs with OpenMP and OpenCL," Languages and Compilers for Parallel Computing, pp. 215–229, 2011.
- [18] O. Fagerlund, "Multi-core programming with OpenCL: performance and portability in a memory bound scenario," June 2010. Available at <http://daim.idi.ntnu.no>.
- [19] J. Stone, D. Gohara, and G. Shi, "OpenCL: A parallel programming standard for heterogeneous computing systems," Computing in science & engineering, vol. 12, no. 3, p. 66, 2010.
- [20] A. Duran, E. Ayguad'e, R. M. Badia, J. Labarta, L. Martinell, X. Martorell and J. Planas, "Ompss: a Proposal for Programming Heterogeneous Multi-Core Architectures," Parallel Processing Letters, vol. 21, no. 2, pp. 173–193, 2011.
- [21] Kazuhiko Komatsu, Katsuto Sato, Yusuke Arai, Kentaro Koyama, Hiroyuki Takizawa, Hiroaki Kobayashi, Evaluating the Performance and Portability of OpenCL Programs, 5th Fifth International Workshop on Automatic Performance Tuning, 2010.
- [22] Yao Zhang, Mark Sinclair II, Andrew A. Chien, Improving Performance Portability in OpenCL Programs, Supercomputing, Lecture Notes in Computer Science Volume 7905, 2013, pp 136-150.
- [23] V. T. Ravi, W. Ma, D. Chiu, and G. Agrawal. Compiler and Runtime Support for Enabling Generalized Reduction Computations on Heterogeneous Parallel Configurations. In ICS 2010, pages 137–146, 2010.
- [24] F. Song and S. Tomov and J. Dongarra. Enabling and Scaling Matrix Computations on Heterogeneous Multi-core and Multi-GPU systems. In ICS 2012, pages 365–376, 2012
- [25] M. D. Linderman, J. D. Collins, H. Wang, and T. H. Y. Meng. Merge: A Programming Model for Heterogeneous Multi-core Systems. In ASPLOS 2008, pages 287–296, 2008.
- [26] S. Tomov, J. Dongarra, and M. Baboulin. Towards Dense Linear Algebra for Hybrid GPU Accelerated Manycore Systems. Parallel Computing, 36(5-6):232–240, 2010.
- [27] D. Grewe and M. F. P. O'Boyle. A Static Task Partitioning Approach for Heterogeneous Systems Using OpenCL. In CC 2011, pages 286–305.
- [28] K. Kofler, I. Grasso, B. Cosenza, and T. Fahringer. An Automatic Input-Sensitive Approach for Heterogeneous Task Partitioning. In ICS 2013, pages 149–160, 2013.
- [29] C.-K. Luk, S. Hong, and H. Kim. Qilin: Exploiting Parallelism on Heterogeneous Multiprocessors with Adaptive Mapping. In MICRO 2009, pages 45–55, 2009.
- [30] A. Snaveley, L. Carrington, N. Wolter, J. Labarta, R.M. Badia, A. Purkayastha: A framework for performance modeling and prediction. In SC 2002.
- [31] S. Hong and H. Kim. An Integrated GPU Power and Performance Model. In ISCA 2010,

pages 280–289, 2010.

- [32] Abdullah Gharaibeh, Lauro Beltrão Costa, Elizeu Santos-Neto, and Matei Ripeanu. 2012. A yoke of oxen and a thousand chickens for heavy lifting graph processing. In Proceedings of the 21st international conference on Parallel architectures and compilation techniques (PACT '12).
- [33] Abdullah Gharaibeh, Lauro Beltrao Costa, Elizeu Santos-Neto, and Matei Ripeanu. 2013. On Graphs, GPUs, and Blind Dating: A Workload to Processor Matchmaking Quest. In Proceedings of the 2013 IEEE 27th International Symposium on Parallel and Distributed Processing (IPDPS '13)
- [34] Thomas Scogland, Barry Rountree, Wu-chun Feng, Bronis R. de Supinski: Heterogeneous Task Scheduling for Accelerated OpenMP. IPDPS 2012
- [35] Sylvain Henry, Alexandre Denis, Denis Barthou, Marie-Christine Counilh, and Raymond Namyst. Toward OpenCL Automatic Multi-Device Support. In Fernando Silva, Ines Dutra, and Vitor Santos Costa, editors, Euro-Par 2014, Porto, Portugal, August 2014. Springer.
- [36] Cédric Augonnet, Samuel Thibault, Raymond Namyst, and Pierre-André Wacrenier. StarPU: A Unified Platform for Task Scheduling on Heterogeneous Multicore Architectures. Concurrency and Computation: Practice and Experience, Special Issue: Euro-Par 2009, 23 :187–198, February 2011.